

Caching Spring Boot Microservices with Hazelcast in Kubernetes

Deployment Guide

By Enes Ozcan

Table of Contents

Caching Spring Boot Microservices with Hazelcast in Kubernetes	3
■ What You'll Learn	3
■ What Is Hazelcast?	3
■ Why Spring Boot?	3
■ Prerequisites	3
■ Getting Started	4
■ Running a Spring Application	4
■ Dockerizing the App.....	5
■ Running the App in a Container.....	6
■ Starting and Preparing Your Cluster for Deployment	6
Validate Kubernetes environment	6
■ Hazelcast Caching Among Kubernetes Pods	8
■ Scaling with Hazelcast	9
■ Testing Microservices Running on Kubernetes.....	10
■ Tearing Down the Environment	11
■ Conclusion	12



Caching Spring Boot Microservices with Hazelcast in Kubernetes

You can see the whole project [here](#) and start building your app in the final directory.

What You'll Learn

In this guide, you will learn how to use Hazelcast distributed caching with Spring Boot and deploy to a local Kubernetes cluster. You will then create a Kubernetes Service which load balances between containers and verify that you can share data between microservices.

The microservice you will deploy is called `hazelcast-spring`. The `hazelcast-spring` microservice simply helps you put data and read it back. The Kubernetes Service will send the request to a different pod each time you initiate the request, and the data will be served by a shared hazelcast cluster between `hazelcast-spring` pods.

You will use a local single-node Kubernetes cluster. However, you can deploy this application on any Kubernetes distributions.

What Is Hazelcast?

Hazelcast is an open source in-memory data grid (IMDG). It provides elastically scalable, distributed in-memory computing, widely recognized as the fastest and most scalable approach to application performance.

Hazelcast is designed to scale up to hundreds and thousands of members. Simply add new members and they will automatically discover the cluster and will linearly increase both memory and processing capacity.

Why Spring Boot?

Spring Boot makes it easy to create stand-alone, production-grade Spring-based applications that you can “just run.” To learn more about Spring Boot, visit this [website](#).

Prerequisites

Before you begin, have the following tools installed:

- You will need Apache Maven to build and run the project.
- You will also need a containerization software for building containers. Kubernetes supports a variety of container types. We will use Docker in this guide. For installation instructions, refer to the [official Docker documentation](#).

For Windows and Mac

- Use Docker Desktop, where a local Kubernetes environment is pre-installed and enabled. Download it from the [official website](#).

Linux

- Use Minikube as a single-node Kubernetes cluster that runs locally in a virtual machine. For Minikube installation instructions, please [visit this page](#).

Getting Started

The fastest way to work through this guide is to clone the Git repository and use the projects that are provided inside:

```
$ git clone https://github.com/enozcan/guide-kubernetes-caching-hazelcast-spring.git
$ cd guide-kubernetes-caching-hazelcast-spring
```

The `initial` directory contains the starting project that you will build upon.

The `final` directory contains the finished project you will build.

Running a Spring Application

The application in the `initial` directory is a basic Spring Boot app having 3 endpoints:

- `/` is the homepage returning "Welcome" string only
- `/put` is the page where key and values can be put on a concurrent hash map.
- `/get` is the page where the values in the map can be obtained by keys.

Build the app using Maven in the `initial` directory:

```
$ > mvn package
```

Run the application:

```
$ > java -jar target/hazelcast-spring-app-0.1.0.jar
```

Now your app is running on `localhost:8080`. You can test by following requests:

```
$ > curl "localhost:8080"
$ > curl "localhost:8080/put?key=key1&value=hazelcast"
$ > curl "localhost:8080/get?key=key1"
```

This part was the introduction of the application. You can stop your application by **CTRL + C**.

Dockerizing the App

To create the Docker image of the application, add following lines into pom.xml file:

```
<!-- add among other properties -->
<properties>
  <docker.image.prefix>springio</docker.image.prefix>
</properties>

<!-- add among other plugins -->
<plugins>
<plugin>
  <groupId>com.spotify</groupId>
  <artifactId>dockerfile-maven-plugin</artifactId>
  <version>${version.dockerfile-maven-plugin}</version>
  <configuration>
    <repository>${docker.image.prefix}/${project.artifactId}</repository>
  </configuration>
</plugin>
</plugins>
```

Then create the Dockerfile under initial directory named “Dockerfile” containing the instructions for creating a Docker image:

```
FROM openjdk:8-jdk-alpine
VOLUME /tmp
ARG JAR_FILE=target/hazelcast-spring-app-0.1.0.jar
ADD ${JAR_FILE} hazelcast-spring-demo.jar
ENTRYPOINT ["java","-Djava.security.egd=file:/dev/./urandom","-jar","/hazelcast-spring-demo.jar"]
```

Before creating the Docker image of the app, first rebuild the app:

```
$ > mvn clean package
```

Then create image:

```
$ > docker build -t hazelcast-spring-demo .
```

Now, the image must be seen among the docker images:

```
$ > docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
hazelcast-spring-demo	latest	5f65a62b0aa0	19 seconds ago	122MB

Running the App in a Container

Now that the Docker image is ready, check if the image runs properly:

```
$ > docker run -p 5000:8080 hazelcast-spring-demo
```

Test the app on the port 5000:

```
$ > curl "localhost:5000"  
$ > curl "localhost:5000/put?key=key1&value=hazelcast"  
$ > curl "localhost:5000/get?key=key1"
```

If you see the same responses as the ones you get when the app is run without the container, that means it's all OK with the image.

To stop the container, get the container ID first:

```
$ > docker ps
```

Then find the application's container ID and stop the container:

```
$ > docker stop [CONTAINER-ID]
```

Starting and Preparing Your Cluster for Deployment

Now that you have a proper Docker image, deploy the app to Kubernetes pods. Start your Kubernetes cluster first.

Windows | Mac

Start your Docker Desktop environment. Make sure "Docker Desktop is running" and "Kubernetes is running" status are updated.

Linux

Run the following from command line:

```
$ > minikube start
```

Validate Kubernetes environment

Next, validate that you have a healthy Kubernetes environment by running the following command from the command line.

```
$ > kubectl get nodes
```

Hazelcast IMDG | Deployment Guide

This command should return a **Ready** status for the master node.

Windows | Mac

You do not need to do any other step.

Linux

Run the following command to configure the Docker CLI to use Minikube's Docker daemon.

After you run this command, you will be able to interact with Minikube's Docker daemon and build new images directly to it from your host machine:

```
$ > eval $(minikube docker-env)
```

After you're sure that a master node is ready, create `kubernetes.yaml` under `initial` directory with the same content in the `final/kubernetes.yaml` file.

This file defines two Kubernetes resources: one `StatefulSet` and one `service`. `StatefulSet` is preferred solution for Hazelcast because it enables controlled scale out/in of your microservices for easy data distribution. To learn more about `StatefulSet`, you can visit Kubernetes [documentation](#).

By default, we create 2 replicas of `hazelcast-spring` microservice behind the `hazelcast-spring- service` which forwards requests to one of the pods available in the `kubernetes` cluster.

`MY_POD_NAME` is an environment variable made available to the pods so that each microservice knows which pod they are in. This is going to be used in this guide in order to show which pod is responding to the HTTP request.

Run the following command to deploy the resources as defined in `kubernetes.yaml`:

```
$ > kubectl apply -f kubernetes.yaml
```

Run the following command to check the status of your pods:

```
$ > kubectl get pods
```

You'll see an output similar to the following if all the pods are healthy and running:

NAME	READY	STATUS	RESTARTS	AGE
Hazelcast-spring-statefulset-0	1/1	Running	0	7s
Hazelcast-spring-statefulset-1	1/1	Running	0	5s

Send request to port `:31000` and see the pods responding.

```
$ > curl localhost:31000
```

And add a value to the map and then get the value:

```
$ > curl "localhost:31000/put?key=key1&value=hazelcast"

{"value":"hazelcast","podName":"hazelcast-spring-statefulset-1"}

$ > while true; do curl localhost:31000/get?key=key1; echo; sleep 2; done

{"value":"hazelcast","podName":"hazelcast-spring-statefulset-1"}
{"value":null,"podName":"hazelcast-spring-statefulset-0"}
{"value":"hazelcast","podName":"hazelcast-spring-statefulset-1"}
{"value":null,"podName":"hazelcast-spring-statefulset-0"}
```

As can be seen, data is not shared between nodes. Here is where Hazelcast comes into action. Kill active pods under `initial` directory by:

```
$ > kubectl delete -f kubernetes.yaml
```

Hazelcast Caching Among Kubernetes Pods

Now we will use Hazelcast caching among the pods. Update the `pom.xml` file by adding those dependencies:

```
<dependency>
  <groupId>com.hazelcast</groupId>
  <artifactId>hazelcast</artifactId>
  <version>${version.hazelcast}</version>
</dependency>

<dependency>
  <groupId>com.hazelcast</groupId>
  <artifactId>hazelcast-kubernetes</artifactId>
  <version>${version.hazelcast-kubernetes}</version>
</dependency>
```

Then modify the `CommandController.java` such that Hazelcast is used in the map. Also add Hazelcast config to `Application.java` file and import Hazelcast libraries as well. Those versions are the ones under `final` directory and can be copied directly into `initial` directory.

Rebuild the app and create new image:

```
$ > mvn clean package
$ > docker build -t hazelcast-spring-demo .
```

Before deploying on kubernetes, create `rbac.yaml` file as in the `final` directory. The role-based access control (RBAC) configuration is used to give access to the Kubernetes Master API from pods which runs microservices. Hazelcast requires read access to autodiscover other Hazelcast members and form Hazelcast cluster.

Hazelcast IMDG | Deployment Guide

Run the following commands to deploy the resources as defined in `kubernetes.yaml` and `rbac.yaml` in the specified order:

```
$ > kubectl apply -f rbac.yaml
$ > kubectl apply -f kubernetes.yaml
```

Run the following command to check the status of your pods:

```
$ > kubectl get pods
```

You should also check if the Hazelcast cluster is formed by checking one of the pod's log files:

```
$ > kubectl logs hazelcast-spring-statefulset-1
```

You must see such a response at the end of the log:

```
Members {size:2, ver:2} [
  Member [10.1.0.52]:5701 - ac54036d-c16f-40ae-9531-93e6f0683cf9 this
  Member [10.1.0.53]:5701 - d963bb82-3842-49fd-a522-82c8543bdb9d
]
```

If it's not seen, wait for pods to be configured and try again.

Now we expect all nodes to give the same value for the same key put on the map via one pod only. Let's try:

```
$ > curl "http://localhost:31000/put?key=key1&value=hazelcast"
{"value":"hazelcast","podName":"hazelcast-spring-statefulset-1"}

$ > while true; do curl localhost:31000/get?key=key1;echo; sleep 2; done
{"value":"hazelcast","podName":"hazelcast-spring-statefulset-1"}
{"value":"hazelcast","podName":"hazelcast-spring-statefulset-1"}
{"value":"hazelcast","podName":"hazelcast-spring-statefulset-0"}
```

As can be seen, the insertion is made on `hazelcast-spring-statefulset-1` but both nodes gives the same value for the key now.

Scaling with Hazelcast

Scale the cluster with one more pod and see that you still retrieve the shared data.

```
$ > kubectl scale statefulset hazelcast-spring-statefulset --replicas=3
```

Run following command to see the latest status of the pods

```
$ > kubectl get pods
```

As you can see, a new pod `hazelcast-spring-statefulset-2` has joined the cluster.

NAME	READY	STATUS	RESTARTS	AGE
Hazelcast-spring-statefulset-0	1/1	Running	0	8m
Hazelcast-spring-statefulset-1	1/1	Running	0	8m
Hazelcast-spring-statefulset-2	1/1	Running	0	30s

Run the following command again to see the output

```
$ > while true;do curl http://localhost:31000/get?key=key1;echo; sleep 2; done
```

```
{"value":"hazelcast","podName":"hazelcast-spring-statefulset-2"}  
{"value":"hazelcast","podName":"hazelcast-spring-statefulset-0"}  
{"value":"hazelcast","podName":"hazelcast-spring-statefulset-2"}  
{"value":"hazelcast","podName":"hazelcast-spring-statefulset-1"}
```

As you can see, `hazelcast-spring-statefulset-2` is returning correct data.

Testing Microservices Running on Kubernetes

Create a testing class under `initial/src/test/java/it/io/spring/guides/hazelcast/` named `HazelcastCachingIT.java`. The contents of the test file is available under the `final` directory.

The test makes sure that the `/put` endpoint is handled by one pod and `/get` methods returns the same data from another Kubernetes pod.

It first puts a `key/value` pair to the `hazelcast-spring` microservice and saves the pod name in the `firstpod` variable. In the second part, the test submits multiple `/get` requests until a different pod name is seen other than the pod which initially handled `/put` request.

In order to run integration tests, you must have running `hazelcast-spring` microservices in a `minikube` environment. As you have gone through all the previous steps, you already have it.

Navigate back to `initial` directory and run the test:

```
$ > mvn -Dtest=HazelcastCachingIT test
```

If the tests pass, you'll see a similar output to the following:

```
[INFO] TESTS
[INFO] -----
[INFO] Running HazelcastCachingIT
10:12:27.087 [Time-limited test] DEBUG org.springframework.web.client.RestTemplate -
HTTP GET http://localhost:31000/put?key=key1&value=hazelcast-spring-guide
10:12:27.175 [Time-limited test] DEBUG org.springframework.web.client.RestTemplate -

Accept=[application/json, application/*+json]
10:12:27.312 [Time-limited test] DEBUG org.springframework.web.client.RestTemplate -
Response 200 OK
...
...
[INFO] Tests run: 1, Failures: 0, Errors: 0, Skipped: 0, Time elapsed: 5.354 s - in
HazelcastCachingIT
[INFO]
[INFO] Results:
[INFO]
[INFO] Tests run: 1, Failures: 0, Errors: 0, Skipped: 0
[INFO]
[INFO] -----
[INFO] BUILD SUCCESS
```

Tearing Down the Environment

When you no longer need your deployed microservices, you can delete all Kubernetes resources by running the `kubectl delete` command: You might need to wait up to 30 seconds as stateful sets kills pods one at a time.

```
$ > kubectl delete -f kubernetes.yaml
```

Windows | Mac

Nothing more needs to be done for Docker Desktop.

Linux

Perform the following steps to return your environment to a clean state.

Point the Docker daemon back to your local machine:

```
$ > eval $(minikube docker-env -u)
```

Stop your Minikube cluster:

```
$ > minikube stop
```

Delete your cluster:

```
$ > minikube delete
```

Conclusion

You have just run a Spring Boot application and created its Docker image. First you ran the app on a container and then deployed it to Kubernetes. You then added Hazelcast caching to the hazelcast-spring, tested with a simple curl command. You also scaled out the microservices and saw that data is shared between microservices. As a last step, you ran integration tests against hazelcast-spring that was deployed in a Kubernetes cluster.



350 Cambridge Ave, Suite 100, Palo Alto, CA 94306 USA
Email: sales@hazelcast.com Phone: +1 (650) 521-5453
Visit us at www.hazelcast.com

Hazelcast, and the Hazelcast, Hazelcast Jet and Hazelcast IMDG logos are trademarks of Hazelcast, Inc. All other trademarks used herein are the property of their respective owners. ©2019 Hazelcast, Inc. All rights reserved.

